

# How to make a (GameBoy) emulator?

Petar Veličković

University of Cambridge  
CAMBRIDGE CODING SUMMER SCHOOL

13 July 2016

# Emulation

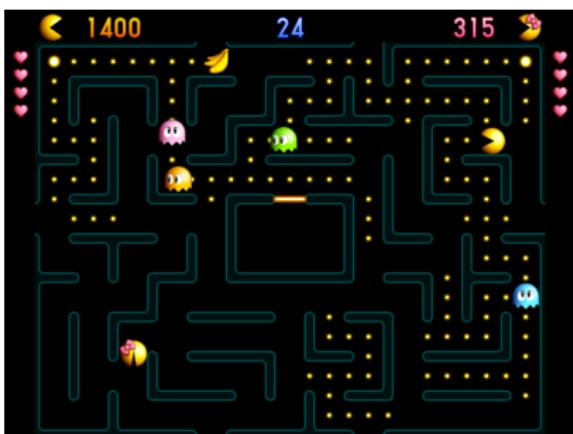
- ▶ **Emulation** is a computer science concept corresponding to *imitation of one computer system on another*.
  - ▶ Contrast with **simulation**, which is just an imitation of a particular feature/functionality (e.g. of a video game).
  - ▶ Emulation can be performed in both **hardware** and **software**; in this talk, we will focus only on software emulators.

## What is emulation?

## Emulator vs simulator



## PSX emulator



## Pac-man simulator

An emulator executes the **original code (ROM)** written for the console, whereas the simulator represents an entirely new game.

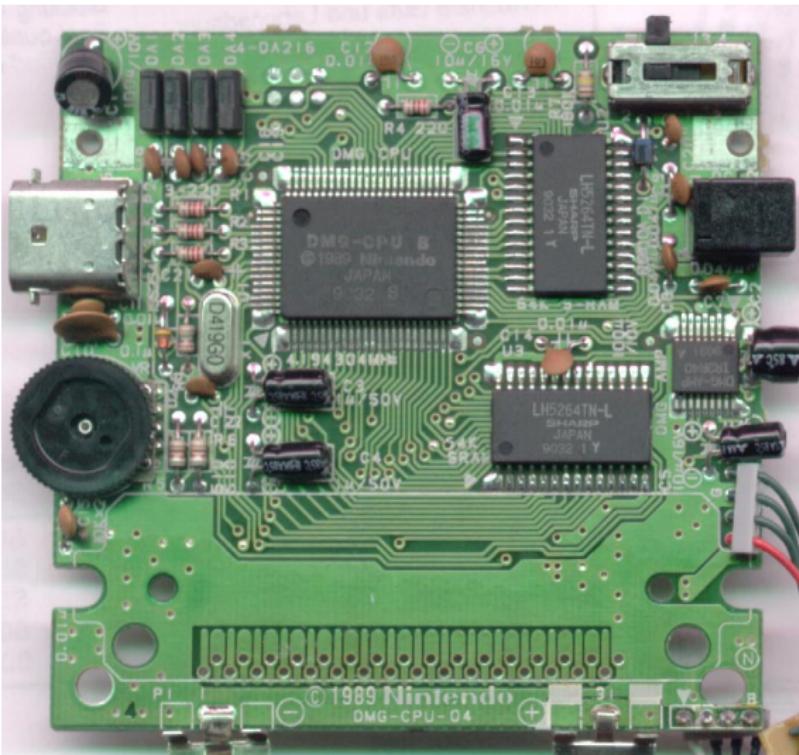
# GameBoy

- ▶ The *GameBoy* is a video gaming console released by Nintendo in 1989.
  - ▶ First notable *handheld* console, in spite of technologically superior competition:
    - ▶ Together with its successor (Game Boy Color), sold **over 119 million units** (by far the most popular handheld console prior to the release of the Nintendo DS).
    - ▶ The Tetris cartridge alone sold 35 million copies!

## Console exterior



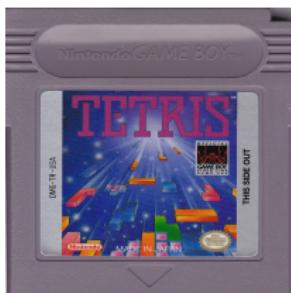
## Under the hood



## Overview of the GameBoy emulator

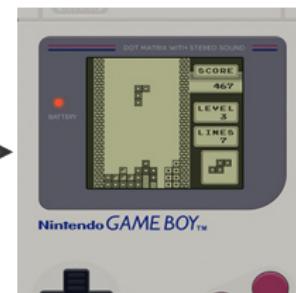
## Today's task

## *Cartridge*



insert

GameBoy



22

۲۷

1100 0011 0000 1100

0000 0010 0000 0000

0000 0000 0000 0000

2000 2010 1111 1111

11111111111111111111

ROM (Game code)

input

## *ROM (Game code)*

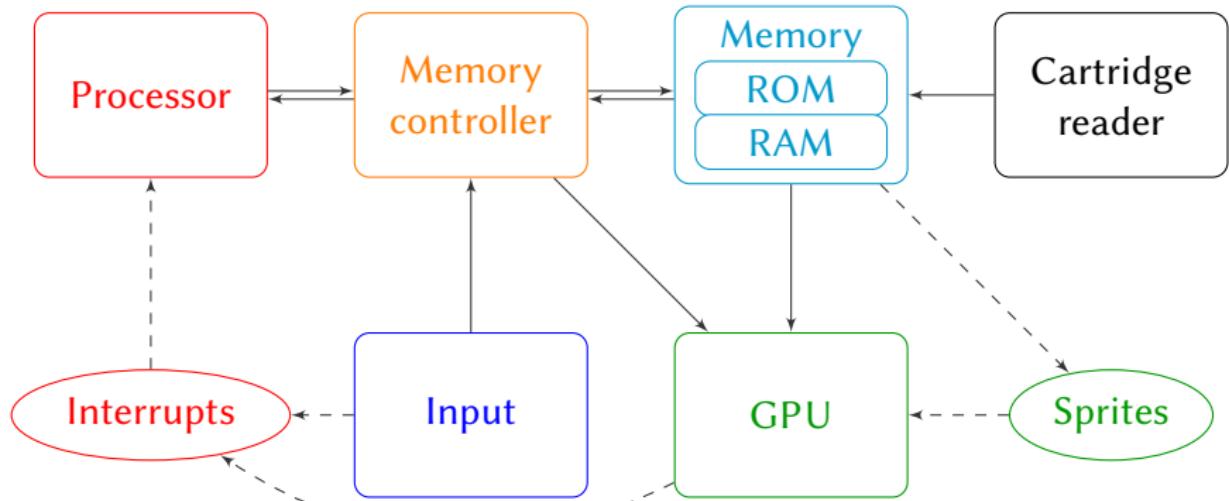


## *Emulator*

# Modular approach

- ▶ We will adopt a *modular* approach in our implementation – the emulator will consist of (nearly) **independent** implementations of all of its subsystems (~ modules)!
- ▶ Why do this?
  - ▶ Debugging becomes much easier because the effects of programmer errors are local to a subsystem.
  - ▶ The modules may be reusable in other emulator projects – the hardware is *standardised*.
  - ▶ Amenable to *object-oriented programming*.

# Overview of the emulator

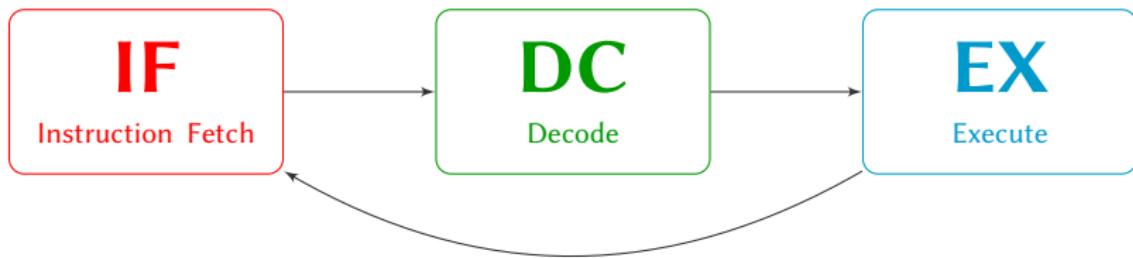


# Processor

Processor

- ▶ Usually the most involved and complex module in the emulator implementation (which is why programmers usually tend to use pre-made CPU implementations).
- ▶ Several approaches are possible—we will focus on writing an **interpreter**, which “follows” the game code and simply simulates every instruction it encounters using the local CPU.

# Interpreter emulator



- ▶ Interpreters usually simulate the *fetch-decode-execute* cycle, present in most contemporary processors:
  - IF** *Instruction Fetch* – fetch the current instruction;
  - DC** *Decode* – determine the type of the instruction;
  - EX** *Execute* – execute the current instruction, updating the state of the processor, memory, etc.
- ▶ This way, we can only emulate systems whose processors are **several orders of magnitude** slower than the native one.

# Instruction Fetch



- ▶ Where is our program???
- ▶ With GameBoy, and most contemporary systems, both the code and data are stored in the same memory ( $\sim$  von Neumann architecture).
- ▶ ⇒ We need methods for reading from and writing to memory, if we'd like to be able to fetch the current instruction.

# Memory interface

- ▶ The GameBoy does all of its computations on the level of a single byte; therefore, it makes sense to expose two methods:

```
ReadByte(address);  
WriteByte(address, value);
```

to the CPU.

- ▶ For now, the processor only needs to know that these functions exist—we will consider their actual implementations later.

# Decode



- ▶ After fetching the current byte of the program, we need to determine the instruction that it represents (and hence whether we need to read additional parameters from memory).
- ▶ This requires us to be well acquainted with the **instruction set architecture** (ISA) of GameBoy's processor (which is very similar to the Zilog Z80 ISA).

# GameBoy architecture

- ▶ Since we have read a single byte of our program as an instruction, this allows us to consider  $2^8 = 256$  distinct instructions.
- ▶ It is necessary to implement an enormous branching (switch-case is an appropriate construct to use), which will select the instruction to execute based on this byte's value.
- ▶ This is easily the most cumbersome part of the implementation...

# Execute



- ▶ Each instruction may affect the processor's internal state (~ **registers**), as well as the state of memory.
- ▶ Furthermore, every instruction requires a certain (predetermined) number of CPU ticks/cycles (**will be very important later!**).
- ▶ The previously mentioned `WriteByte` method may be used for effects on memory.
- ▶ In order to implement all these instructions, we need to implement and maintain the processor's internal state as well.

# Processor state

The CPU's current state is fully specified by the following registers:

PC *Program Counter* — the memory address of the next instruction to be fetched;

SP *Stack Pointer* — the memory address of the top of the stack (omitted from this talk);

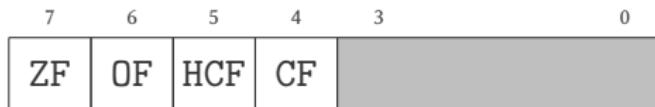
F *Flags*;

A-E,H,L General purpose registers.

- We will also maintain the number of elapsed ticks since the start of execution.

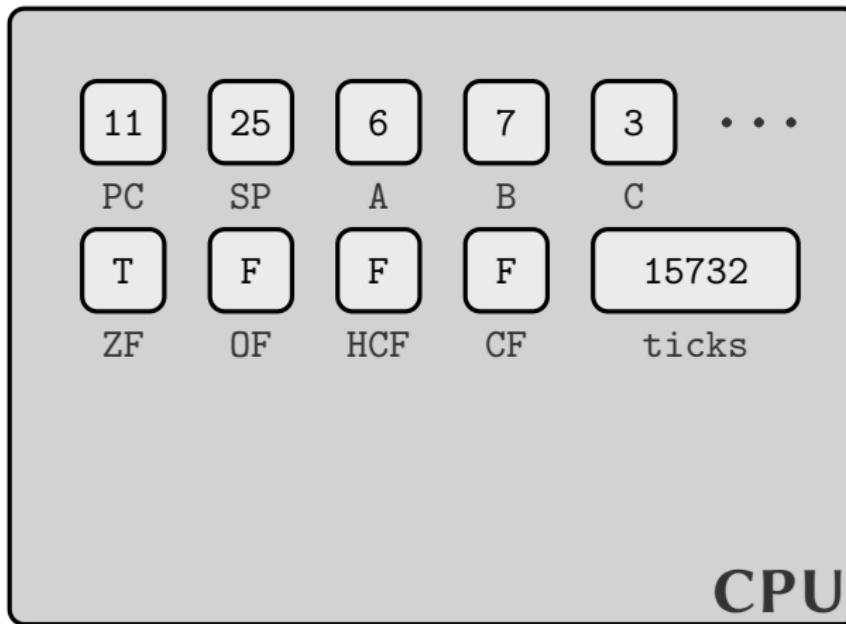
# Flags

- ▶ Flags are logical (true/false) variables that allow the processor to determine the effects of the previously executed instruction.
- ▶ With GameBoy, the flags register looks like this:



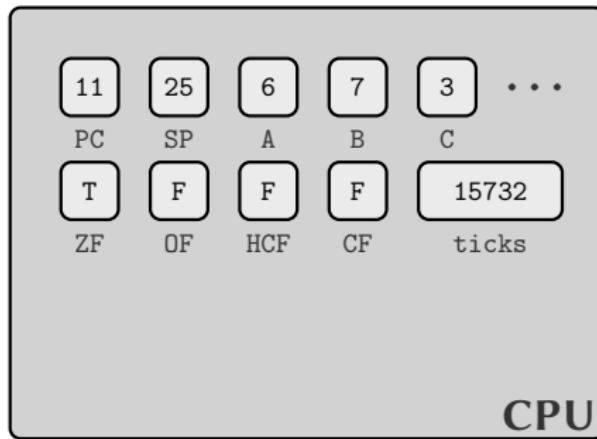
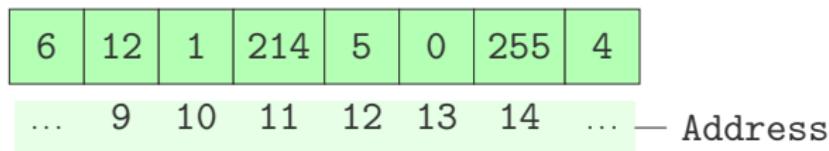
- ▶ The aforementioned flags are:
  - ZF *Zero Flag* — set if the result of the operation is zero;
  - OF *Operation Flag* — set if the operation was subtraction;
  - HCF *Half-Carry Flag* — set if there was an overflow in the lower half of the result (omitted);
  - CF *Carry Flag* — set if there was an overflow in the result.

## CPU state, *cont'd*

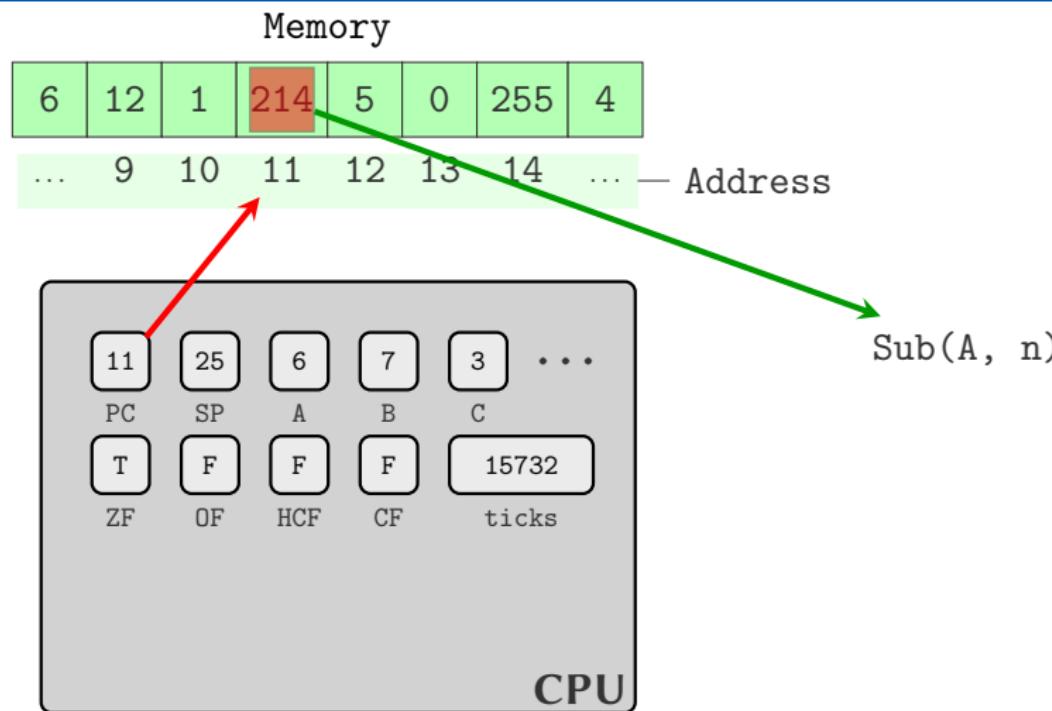


## Example: Sub(A, 5)

Memory



## Example: Sub(A, 5)



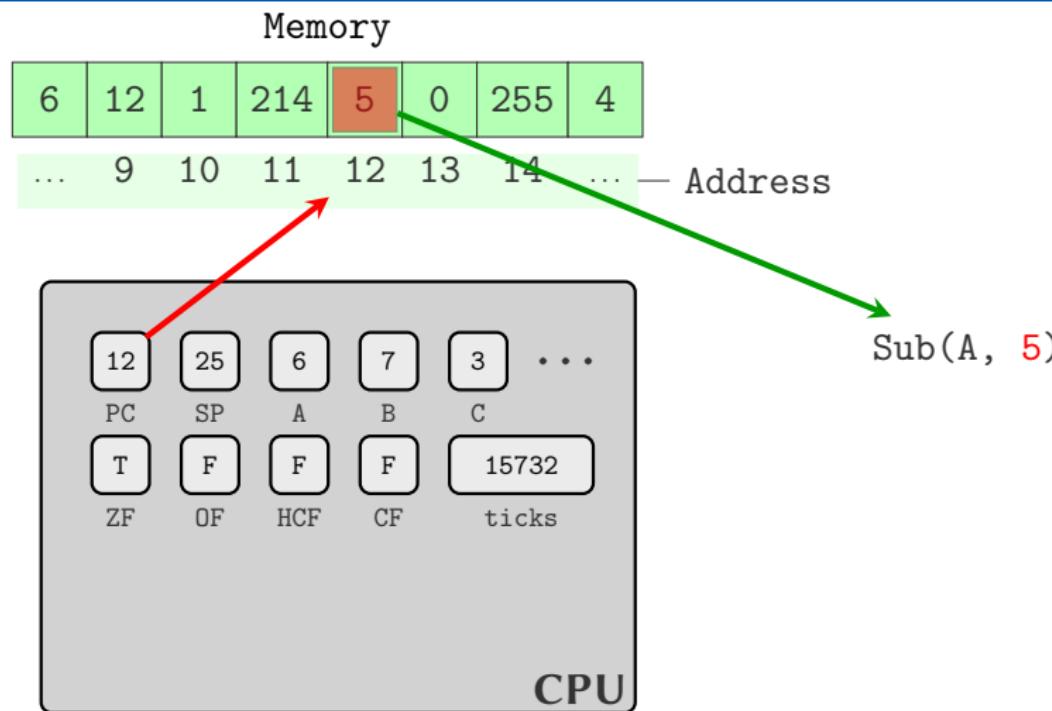
Introduction  
oooooooooooo

Essential emulator  
oooooooooooooooo●oooooooooooo

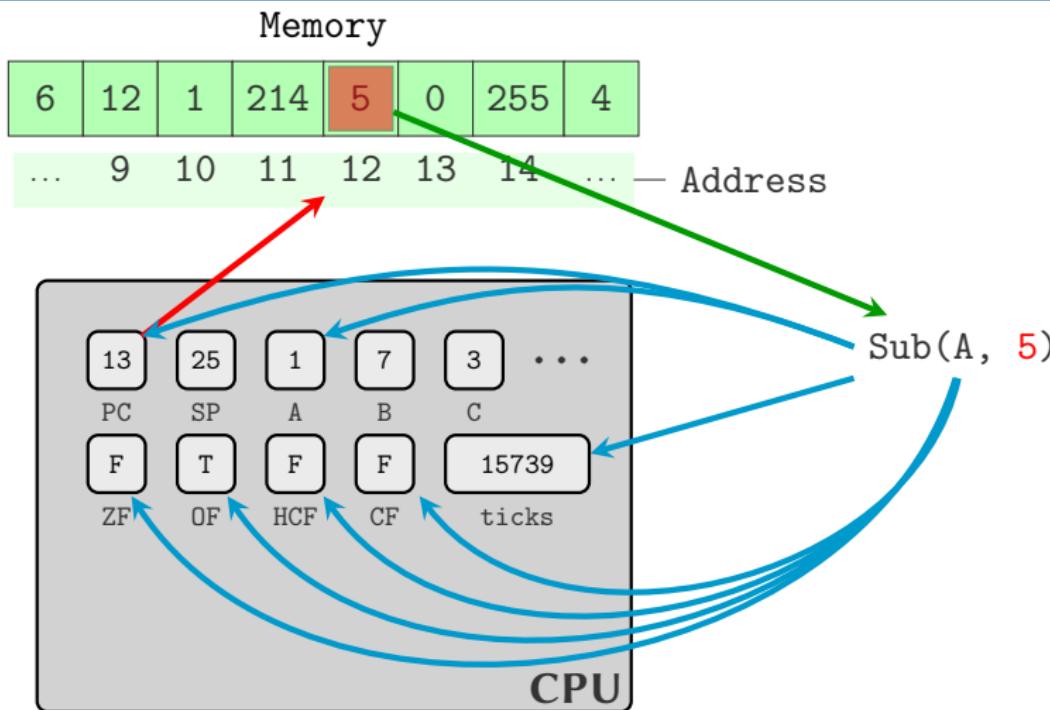
Emulator extensions  
oooooooooooo

Processor

## Example: Sub(A, 5)



## Example: Sub(A, 5)



# Memory



- ▶ It is necessary to correctly implement the `ReadByte` and `WriteByte` methods: we need to emulate the **memory controller**, i.e. the memory management unit (*MMU*).

# Naïve implementation

- ▶ PC has 16 bits, and therefore the GameBoy memory system supports  $2^{16} = 65536$  addresses.
- ▶ The naïve implementation is very simple...

```
1 int memo[65536];
2 public int ReadByte(int address)
3 {
4     return memo[address];
5 }
6 public void WriteByte(int address, int value)
7 {
8     memo[address] = value;
9 }
```

- ▶ Can you think of some reasons why this is not so easy?

# Gotchas

- ▶ Generally, *logical address*  $\neq$  *physical address*:
  - ▶ Not all writes should succeed (ROM...);
  - ▶ Not all addresses represent main memory (e.g. a special address allows the CPU to read keypad inputs...);
  - ▶ Two different addresses could be mapped to the same physical location.
  - ▶ ...
- ▶ With GameBoy, **all of the above hold!**

# GPU

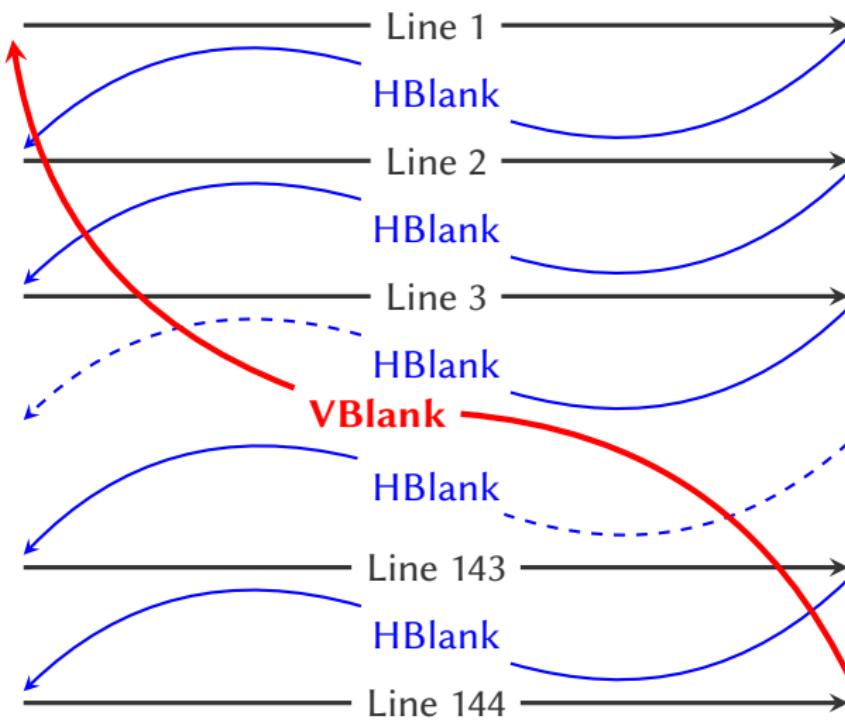
GPU

- ▶ So far, we have managed to design an emulator that can, step by step, execute any GameBoy program.
- ▶ Now we are interested in seeing the results of our work on a graphical output.

# GameBoy graphics

- ▶ GameBoy's graphical output is an LCD with a  $160 \times 144$  pixel resolution  $\implies$  we can maintain a colour matrix of such size, and display it in our emulator.
- ▶ The values in this matrix have the following format:  
`0xAARRGGBB` (Alpha–Red–Green–Blue)
- ▶ GameBoy supports four shades of gray, so the matrix may contain only these four values:
  - ▶ `0xFF000000` (BLACK)
  - ▶ `0xFF555555` (DARK\_GRAY)
  - ▶ `0xFFAAAAAA` (LIGHT\_GRAY)
  - ▶ `0xFFFFFFFF` (WHITE)

# Rendering a single frame ( $\sim$ CRT)



# When do we draw a new frame?

- ▶ Every stage in the previous figure takes a specific amount of time:

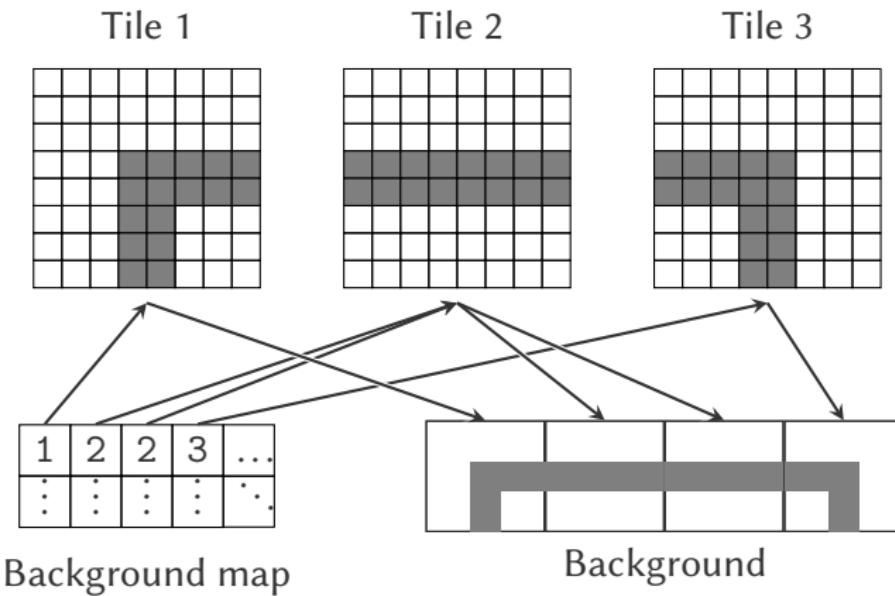
Phase	Ticks
Line (background)	172
Line (sprites)	80
Horizontal blank	204
<i>Single line</i>	456
Vertical blank	4560
<i>Entire frame</i>	70224

- ▶ How could we keep track of *when* to trigger each phase?

# Tiling system

- ▶ GameBoy doesn't have enough memory to keep track of the entire matrix of pixels; instead, we use a *tiling* system.
- ▶ VRAM contains “tiles” of size  $8 \times 8$ , and the background is represented as a  $32 \times 32$  matrix of pointers to the tiles.
- ▶ Observe that this gives us a total resolution of  $256 \times 256$ , yet GameBoy's LCD is  $160 \times 144$ !
  - ▶ Two special-purpose registers, Scroll-X and Scroll-Y, contain the coordinates of the point on the background that will be located in the upper-left corner of the LCD.

## Tiling system, *cont'd*



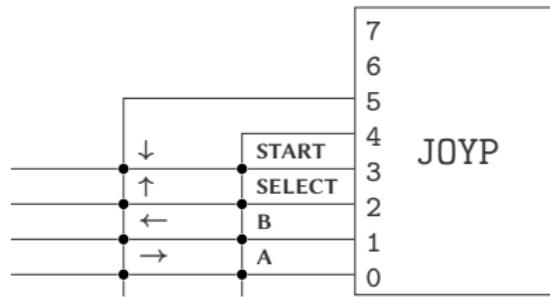
# Input

Input

- ▶ With the GPU implemented, the emulator is able to display graphical output, but the user/player still has no control over the game.
- ▶ Next step: emulating user input.

# Input

- ▶ Input is processed using a special-purpose JOYP register, which the processor may access at address 65280. The wires leading to that register are connected in the following way:



- ▶ The processor first writes 0x10 or 0x20 to the register, so it can activate one of two columns; then, after several ticks, it can read the least significant 4 bits in order to infer which buttons were pressed.

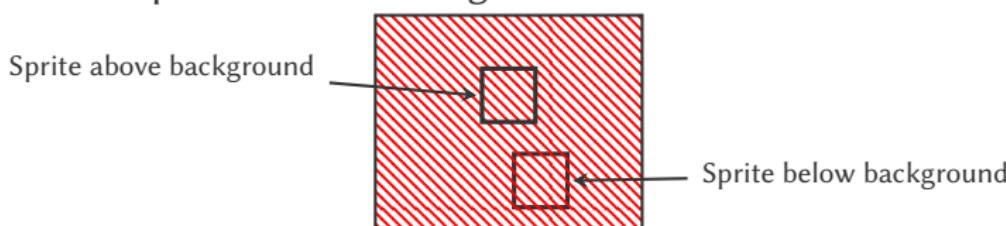
# Sprites

## Sprites

- ▶ With the input processed, we can play some games – but we usually have to play blindly (there is no way the player can tell what his/her position is).
- ▶ This issue (and many others) is solved by using **sprites** – tiles which may be drawn and moved independent of the background.

# Sprites: additional details

- ▶ Sprites use exactly the same type of  $8 \times 8$  tiles as the background.
- ▶ Specific parameters, maintained for every sprite in the OAM:
  - ▶  $(x, y)$  coordinates of the upper-left corner;
  - ▶ *priority*: is it in the foreground or behind the background;
  - ▶ *flip*: horizontal or vertical;
  - ▶ *size*: all sprites can be resized to  $8 \times 16$  at the same time.
- ▶ The frame rendering algorithm: first draw the background, then infer (based on colours and sprite priority) whether to draw the sprite over the background.

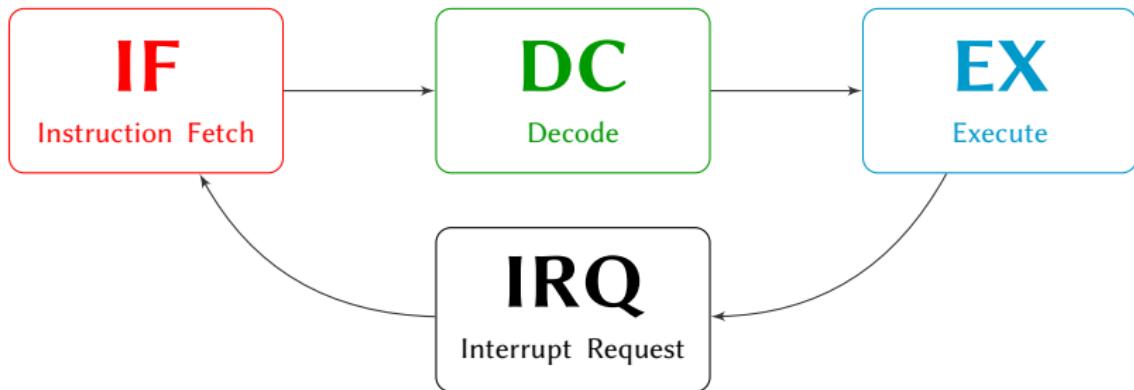


# Interrupts

## Interrupts

- ▶ With sprite support, some simpler games work exactly as expected.
- ▶ Most games require a very important detail: notifying the processor **when** it may render the next frame.
- ▶ The most convenient time for changing the frame is **during VBlank**, because during the line change, the GPU doesn't access either VRAM nor OAM.
- ▶ The best (in terms of resource management and requirements) method for notifying the processor about an important event such as entering VBlank is using **interrupts**.

# Fetch-decode-execute cycle, revisited



- ▶ After every executed instruction, the CPU needs to check whether an interrupt was raised; if an interrupt is detected, the CPU:
  - ▶ Remembers its current state;
  - ▶ Jumps (sets the PC) to the address of the *interrupt handler*.
  - ▶ Executes the interrupt handler code, which ends with a special RETI (*Return From Interrupt*) instruction, which restores state.

# Interrupts: GameBoy specifics

- ▶ The processor has access to the *Interrupt Enable (IE)* registers, through which it can specify which interrupts (out of 5 in total) it is willing to handle at this point.
  - ▶ If multiple interrupts are detected, the one with the highest *priority* is handled first.
- ▶ The processor also has an *Interrupt Master Enable (IME)* switch, which allows it to completely deactivate interrupt handling (this is done during e.g. processing an interrupt, because we should not handle two of them at the same time).

# GameBoy interrupts

Priority	Interrupt	Address of the handler
0	<i>VBlank</i>	0x0040
1	<i>LCD Status</i>	0x0048
2	<i>Timer</i>	0x0050
3	<i>Serial</i>	0x0058
4	<i>Joypad</i>	0x0060

# We're done!

- ▶ Even with only implementing the VBlank interrupt, our emulator is capable of running *Tetris*. :)
- ▶ All details about GameBoy's functionalities, as well as the Game Boy Color console, may be found on the following (extremely useful) reference page:  
<http://problemkaputt.de/pandocs.htm>
- ▶ For any further queries, feel free to contact me:  
`pv273@cam.ac.uk`
- ▶ **Thank you!**